

Group Description

Group No: 25 Based on Spreadsheet

Group Name: ABRACADABRA

Team Member Details~

- 1. Saad Ahmed Pathan (22114077)
- 2. Samio Ayman (22082403)
- 3. Nur Shaheila Ashriza Binti Mohd Saupi (22001745)
- 4. Nurina Humaira Binti Mohd Romzan (22002204)
- 5. Nur Aina Batrisyia Binti Zakaria (23005013)
- 6. Siti Hajar Binti Mohd Nor Azman (22002035)

1. Project Designing

Electricity is essential for economic and social development, enabling nations to achieve higher living standards.

In today's world, effective planning and operation of electricity production, revenue generation from production, and energy consumption are imperative. Understanding how energy generates revenue and is utilized by consumers is crucial for better management. This presents an opportunity to develop a supervised machine learning model to forecast future electricity revenues.

- 1. Initial Phase: We brainstormed the problem and potential approaches to solve it using machine learning concepts. Then, we designed the workflow of our project.
- 2. Data Mining: We extracted a dataset from Data.gov, covering data from 2015 to 2022. The dataset includes revenue, units sold, and the average number of customers, categorized by customer class for each electric utility operating in Iowa, USA.
- 3. Data Preprocessing: We understood the data and identified some null values in the dataset, receiving a detailed description of the characteristics involved.
- 4. Feature Discussion: We discussed and renamed features for better readability and understanding, facilitating a smoother data environment.
- 5. Exploratory Data Analysis (EDA) and Visualization: EDA and visualization provided concise knowledge of the link between features and the label (the dependent variable). The heatmap was used to understand the association between independent variables, helping to choose important features. Selecting the right elements to improve accuracy was challenging.
- 6. Feature Selection: We decided to use PCA for feature selection, ultimately choosing PC1 as the feature for our project.
- 7. Model Training and Assessment: We employed Linear Regression, Random Forest Regression, Neural Network Regression, Decision Tree, and XGBoost techniques. After comparing numerous metrics, we determined that the Random Forest Regressor produced the best results.
- 8. Model Explainability: We used a bar chart to compare the performance of all five models, assisting in selecting the best one. The Random Forest Regressor emerged as the best model for our dataset.
- 9. Conclusion: We summarized our project, from model selection and evaluation to finding the most suitable model for our dataset. We also highlighted key findings from each model with their respective values.

Problem Statement

The goal is to develop a machine learning model capable of accurately forecasting electricity revenues based on the provided features. This model is valuable for utility companies, energy firms, and policymakers who need to optimize electricity consumption, reduce costs, and minimize the environmental impact of energy usage.

Specifically, the model should reliably predict electricity revenues by considering various factors influencing energy consumption, such as consumer types and the number of consumers. This can help utility companies, building managers, and energy firms identify patterns and trends in energy consumption, enabling them to make informed energy decisions. Policymakers can also use this data to create regulations and incentives that promote energy efficiency and sustainability.

2. Data Mining

The dataset used for this project is acquired from the website Data.gov. Data.gov is a comprehensive and open data portal maintained by the United States government. It serves as a centralized repository for accessing a wide range of government datasets, providing the public, researchers, and developers with valuable information for analysis, innovation, and transparency.

The dataset titled "**Electric Utilities Revenue, Units Sold, and Customers by Year**" covers data from 2015 to 2022, detailing the revenue, units sold, and average number of customers categorized by customer class for each electric utility operating in the state of Iowa, USA. This publicly accessible dataset aims to provide insights into the performance and customer base of electric utilities in Iowa. However, no specific license information is provided for this dataset.

Columns Description

- 1. Reporting Year
- 2. Company Number & Year
- 3. Type of Utility
- 4. Utility
- 5. Operating Revenues - Residential Sales
- 6. Operating Revenues - Commercial & Industrial Sales
- 7. Operating Revenues - Sales for Resale
- 8. Operating Revenues - All Other Sales
- 9. MWh Sold - Residential

10. MWh Sold - Commercial & Industrial
11. MWh Sold - Sales for Resale
12. MWh Sold - All Other
13. Average No. of Customers - Residential
14. Average No. of Customers - Commercial & Industrial
15. Average No. of Customers - Sales for Resale
16. Average No. of Customers - All Other

Dataset Source Link

<https://catalog.data.gov/dataset/electric-utilities-revenue-units-sold-and-customers-by-year>

3. Data Preprocessing

```
1 # Line Wrapping in Collaboratory Google results
2 from IPython.display import HTML, display
3
4 def set_css():
5     display(HTML('''
6     <style>
7         pre {
8             white-space: pre-wrap;
9         }
10    </style>
11    '''))
12 get_ipython().events.register('pre_run_cell', set_css)
```



Check for missing values, outliers, and inconsistencies in the dataset and handle them appropriately. Missing values can be imputed or dropped based on the extent of missingness and their impact on the analysis.

```
1 # Import Libraries for analysis and visualisation
2 import pandas as pd
3 import numpy as np
4 import matplotlib.pyplot as plt
5 import seaborn as sns
6 import missingno as msno
7 %matplotlib inline
```



```
1 # To import datetime library
2 from datetime import datetime
3 import datetime as dt
4
5 # Library of warnings would assist in ignoring warnings issued
6 import warnings
7 warnings.filterwarnings('ignore')
8
9 # Import necessary statistical libraries
10 import scipy.stats as stats
11 import statsmodels.api as sm
12 from scipy.stats import norm
```



```
1 # Import libraries for ML-Model
2 from sklearn.model_selection import train_test_split
3 from sklearn.metrics import mean_absolute_error, r2_score, mean_squared_error
4 from sklearn.preprocessing import StandardScaler, LabelEncoder
5 from sklearn.model_selection import GridSearchCV
6
7 # Libraries for save the model
8 import pickle
```



```
1 # Mount Google Drive to access the dataset
2 from google.colab import drive
3 drive.mount('/content/drive')
4
5 Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).
```



```
1 # Load the dataset
2 file_path = '/content/drive/MyDrive/Machine Learning Project/electricity_consumption_data.csv'
3
4 df = pd.read_csv(file_path)
```



```
1 # Display the shape of the data
2 df.shape
```



(1467, 15)

```
1 # Display the first few rows to understand the data
2 print(df.head())
```

```

RY      ToU      U      \
0  2015  Municipal Electric  Auburn
1  2015  Investor Owned Electric  Interstate Power and Light Company
2  2015  Investor Owned Electric  MidAmerican Energy Company
3  2015  Investor Owned Electric  Amana Society Service Company
4  2015  Municipal Electric  Bloomfield

ORoRS      ORoCIS      ORoSR      ORoAOS      ASforR      ASforCI      \
0  232546.0      0.0      0.0      12494.0      1646.39      0.00
1  521115322.0      846803408.0      74454294.0      11626180.0      3661188.00      10691600.00
2  535517779.0      792025980.0      158876153.0      96997319.0      5490294.00      14032377.00
3  894691.0      6971861.0      0.0      76411.0      7095.91      84452.11
4  1647288.0      1503769.0      0.0      86996.0      11501.42      14118.17

ASforSR      ASforAO      ANoCR      ANoCCI      ANoCSR      ANoCAO
0  0.0      103.25      175      0      0.0      10
1  1779026.0      53530.00      408969      78348      6.0      1040
2  7907806.0      1399758.00      568142      81892      5.0      12764
3  0.0      853.85      714      302      0.0      24
4  0.0      1328.69      1123      268      0.0      1
```

```
1 df.head(5)
```

```

RY      ToU      U      ORoRS      ORoCIS      ORoSR      ORoAOS      ASforR      ASforCI      ASforSR      ASforAO      ANoCR      A
0  2015  Municipal Electric  Auburn  232546.0      0.0      0.0      12494.0      1646.39      0.00      0.0      103.25      175
1  2015  Investor Owned Electric  Interstate Power and Light Company  521115322.0      846803408.0      74454294.0      11626180.0      3661188.00      10691600.00      1779026.0      53530.00      408969
2  2015  Investor Owned Electric  MidAmerican Energy Company  535517779.0      792025980.0      158876153.0      96997319.0      5490294.00      14032377.00      7907806.0      1399758.00      568142
```

Next steps: [Generate code with df](#) [View recommended plots](#)

```
1 df.iloc[745 : 751]
```

```

RY      ToU      U      ORoRS      ORoCIS      ORoSR      ORoAOS      ASforR      ASforCI      ASforSR      ASforAO      ANoCR      ANoCCI      ANoCSR
745  2019  Municipal Electric  Cascade  846082.0      773424.0      0.0      216231.0      7694.0      8566.0      0.0      2281.0      949      162      0.0
746  2019  Municipal Electric  Cedar Falls  16139447.0      18557187.0      7804705.0      3410151.0      170532.0      259510.0      318313.0      55319.0      17128      2251      NaN
747  2019  Municipal Electric  Coon Rapids  753173.0      433337.0      819096.0      282901.0      7272.0      3735.0      10971.0      3153.0      606      139      1.0
748  2019  Municipal Electric  Coming  799057.0      466112.0      0.0      567172.0      8467.0      4685.0      0.0      6218.0      745      218      0.0
```

```
1 df.tail(5)
```

```

RY      ToU      U      ORoRS      ORoCIS      ORoSR      ORoAOS      ASforR      ASforCI      ASforSR      ASforAO      ANoCR      ANoCCI      ANoCSR
1462  2022  Distribution Cooperative  T. I. P. Rural Electric Cooperative  10635839.00      5744518.00      0.00      113722.00      83301.0      65148.0      0.0      41.0      6160      334      0
1463  2022  Distribution Cooperative  The Calhoun County Electric Coop. Assn.  4435962.89      799995.49      519366.76      17749.09      31460.0      6956.0      6268.0      0.0      1671      24      2
1464  2022  Distribution Cooperative  United Electric Cooperative, Inc.  1021956.00      111639.00      0.00      0.00      6548.0      930.0      0.0      0.0      435      18      0
```

```
1 df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1467 entries, 0 to 1466
Data columns (total 15 columns):
#   Column      Non-Null Count  Dtype
---  -
0   RY          1467 non-null   int64
1   ToU         1467 non-null   object
2   U           1467 non-null   object
3   ORoRS       1467 non-null   float64
4   ORoCIS      1467 non-null   float64
5   ORoSR       1467 non-null   float64
6   ORoAOS      1467 non-null   float64
7   ASforR      1467 non-null   float64
8   ASforCI     1467 non-null   float64
9   ASforSR     1467 non-null   float64
10  ASforAO     1467 non-null   float64
11  ANoCR       1467 non-null   int64
12  ANoCCI      1467 non-null   int64
13  ANoCSR      1452 non-null   float64
14  ANoCAO      1467 non-null   int64
dtypes: float64(9), int64(4), object(2)
memory usage: 172.0+ KB
```

```
1 # Determine the datatype of Each Column
2 df.dtypes
```

```

RY          int64
ToU         object
U           object
ORoRS       float64
ORoCIS      float64
ORoSR       float64
ORoAOS      float64
ASforR      float64
ASforCI     float64
ASforSR     float64
ASforAO     float64
ANoCR       int64
ANoCCI      int64
ANoCSR      float64
ANoCAO      int64
dtype: object

```

```

1 # Get a statistical summary to check for outliers
2 print(df.describe())

```

```

count  1467.000000  1.467000e+03  1.467000e+03  1.467000e+03  1.467000e+03  \
mean    2018.476483  9.714942e+06  1.404701e+07  3.042812e+06  9.077429e+05
std       2.292097  6.231676e+07  1.036788e+08  2.496384e+07  8.412189e+06
min      2015.000000  0.000000e+00  0.000000e+00  -4.205410e+05  -4.003299e+06
25%     2016.000000  4.169225e+05  2.219630e+05  0.000000e+00  1.498581e+04
50%     2018.000000  9.132270e+05  6.962770e+05  0.000000e+00  6.812389e+04
75%     2020.000000  3.095575e+06  3.372182e+06  0.000000e+00  2.033665e+05
max      2022.000000  6.924891e+08  1.313607e+09  5.470338e+08  1.821175e+08

count  1.467000e+03  1.467000e+03  1.467000e+03  1.467000e+03  1467.000000  \
mean    7.834463e+04  1.868676e+05  9.624706e+04  1.006353e+04  7564.516019
std     5.082949e+05  1.447104e+06  9.128795e+05  1.033581e+05  52795.554431
min     0.000000e+00  0.000000e+00  -6.379000e+03  0.000000e+00  0.000000
25%     3.491000e+03  2.063500e+03  0.000000e+00  5.360500e+01  389.000000
50%     7.926000e+03  6.619000e+03  0.000000e+00  5.058500e+02  758.000000
75%     2.723100e+04  3.644935e+04  0.000000e+00  1.904560e+03  2320.000000
max     6.292473e+06  1.926616e+07  1.753902e+07  1.465272e+06  618886.000000

count  1467.000000  1452.000000  1467.000000
mean    1249.921609  0.832645  108.993865
std     8856.351262  17.361609  955.045804
min      0.000000  0.000000  0.000000
25%      56.000000  0.000000  2.000000
50%     132.000000  0.000000  15.000000
75%     320.000000  0.000000  39.000000
max    90065.000000  660.000000  13186.000000

```

```

1 # Get duplicates count for each unique row
2 dup_count = len(df)-len(df.drop_duplicates())

```

```

1 # There is no duplicate values in the dataframe
2 dup_count1 = df[df.duplicated()].shape
3 dup_count1

```

```
(0, 15)
```

```


1 # Find the missing values of each column
2 null_values = df.isnull().sum()

```

```

1 # Visualizing the missing values
2 plt.figure(figsize=(10,10))
3 sns.displot(
4     data=df.isna().melt(value_name="missing"),
5     y="variable",
6     hue="missing",
7     multiple="fill",
8     aspect=1.25
9 )
10 plt.savefig("visualizing_missing_data_with_barplot_Seaborn_distplot.png", dpi=100)

```



```

1 # Remove all rows with missing data
2 data = df.dropna()
3 data.isna().sum()

```

```
RY      0
ToU     0
U       0
ORoRS   0
ORoCIS   0
ORoSR   0
ORoAOS   0
ASforR   0
ASforCI  0
ASforSR  0
ASforAO  0
ANoCR    0
ANoCCI   0
ANoCSR   0
ANoCAO   0
dtype: int64
```

4. Variable Description

RY - Reporting Year

ToU - Type of Utility

U - Utility

ORoRS - Operating Revenues of Residential Sales

ORoCIS - Operating Revenues of Commercial & Industrial Sales

ORoSR - Operating Revenues of Sales for Resale

ORoAOS - Operating Revenues of All Other Sales

ASforR - Amount Sold for Residential in MWh

ASforCI - Amount Sold for Commercial & Industrial in MWh

ASforSR - Amount Sold for Sales for Resale in MWh

ASforAO - Amount Sold for All Other in MWh

ANoCR - Average No. of Customers in Residential

ANoCCI - Average No. of Customers in Commercial & Industrial

ANoCSR - Average No. of Customers in Sales for Resale

ANoCAO - Average No. of Customers in All Other

```
1 # Show all columns
2 df.columns
```

```
Index(['RY', 'ToU', 'U', 'ORoRS', 'ORoCIS', 'ORoSR', 'ORoAOS', 'ASforR',
       'ASforCI', 'ASforSR', 'ASforAO', 'ANoCR', 'ANoCCI', 'ANoCSR', 'ANoCAO'],
      dtype='object')
```

```
1 df_energy = df.copy()
```

```
1 # Convert to DataFrame
2 df_energy = pd.DataFrame(data)
3
4 # Apply One-Hot Encoding
5 df_energy = pd.get_dummies(df_energy, columns=['ToU', 'U'])
6
7 print("DataFrame after One-Hot Encoding:")
8 print(df_energy)
```

```

DataFrame after One-Hot Encoding:
  RY      ORoRS      ORoCIS      ORoSR      ORoAOS      ASforR \
0  2015  2.325460e+05  0.000000e+00  0.000000e+00  12494.00  1646.39
1  2015  5.211153e+08  8.468034e+08  7.445429e+07  11626180.00  3661188.00
2  2015  5.355178e+08  7.920260e+08  1.588762e+08  96997319.00  5490294.00
3  2015  8.946910e+05  6.971861e+06  0.000000e+00  76411.00  7095.91
4  2015  1.647288e+06  1.503769e+06  0.000000e+00  86996.00  11501.42
...    ...      ...      ...      ...      ...      ...
1462  2022  1.063584e+07  5.744518e+06  0.000000e+00  113722.00  83301.00
1463  2022  4.435963e+06  7.999955e+05  5.193668e+05  17749.09  31460.00
1464  2022  1.021956e+06  1.116390e+05  0.000000e+00  0.00  6548.00
1465  2022  9.256408e+06  4.986241e+06  0.000000e+00  129131.00  67410.00
1466  2022  8.264596e+06  2.162619e+06  0.000000e+00  19479.00  57468.00

      ASforCI      ASforSR      ASforAO      ANoCR      ...      U_Westfield \
0      0.00      0.0      103.25      175      ...      False
1  10691600.00  1779026.0      53530.00  408969      ...      False
2  14032377.00  7907806.0      1399758.00  568142      ...      False
3      84452.11      0.0      853.85      714      ...      False
4      14118.17      0.0      1328.69      1123      ...      False
...    ...      ...      ...      ...      ...      ...
1462  65148.00      0.0      41.00      6160      ...      False
1463  6956.00      6268.0      0.00      1671      ...      False
1464  930.00      0.0      0.00      435      ...      False
1465  37252.00      0.0      601.00      3930      ...      False
1466  19886.00      0.0      24.00      3307      ...      False

      U_Whittemore  U_Wilton  U_Wilton  U_Winterset  U_Winterset \
0      False      False      False      False      False
1      False      False      False      False      False
2      False      False      False      False      False
3      False      False      False      False      False
4      False      False      False      False      False
...    ...      ...      ...      ...      ...
1462      False      False      False      False      False
1463      False      False      False      False      False
1464      False      False      False      False      False
1465      False      False      False      False      False
1466      False      False      False      False      False

      U_Woodbine  U_Woodbine  U_Woodbury County Rural Electric Cooperative \
0      False      False      False
1      False      False      False
2      False      False      False
3      False      False      False
4      False      False      False
...    ...      ...      ...
1462      False      False      False
1463      False      False      False
1464      False      False      False
1465      False      False      False
1466      False      False      True

      U_Woolstock
0      False
1      False
2      False
3      False
4      False
...    ...
1462      False
1463      False
1464      False
1465      False
1466      False

[1452 rows x 239 columns]

```

```

1 # # Rename the columns
2 # df_rename = df.copy()
3 # df_rename.rename(columns={'RY': 'reporting_year', 'ToU': 'utility_type', 'U': 'utility', 'ORoRS ': 'residential_revenues', 'ORoCIS': 'commercial_revenues',
4 # 'ORoSR': 'resale_revenues', 'ORoAOS': 'other_revenues', 'ASforR ': 'residential_sales', 'ASforCI': 'commercial_sales', 'ASforSR': 'resale_sales', 'ASforAO': 'other_sales'
5 # , 'ANoCR': 'residential_customers', 'ANoCCI': 'commercial_customers', 'ANoCSR': 'resale_customers', 'ANoCAO': 'other_customers'}, inplace = True)

```

```

1 df_rename = df.copy()
2 df_rename.rename(columns={
3     'Reporting Year': 'reporting_year',
4     'Company Number & Year': 'company_number_year',
5     'Type of Utility': 'utility_type',
6     'Utility': 'utility',
7     'Operating Revenues - Residential Sales': 'residential_revenues',
8     'Operating Revenues - Commercial & Industrial Sales': 'commercial_revenues',
9     'Operating Revenues - Sales for Resale': 'resale_revenues',
10    'Operating Revenues - All Other Sales ': 'other_revenues',
11    'MWh Sold - Residential': 'residential_sales',
12    'MWh Sold - Commercial & Industrial': 'commercial_sales',
13    'MWh Sold - Sales for Resale': 'resale_sales',
14    'MWh Sold - All Other': 'other_sales',
15    'Average No. of Customers - Residential': 'residential_customers',
16    'Average No. of Customers - Commercial & Industrial': 'commercial_customers',
17    'Average No. of Customers - Sales for Resale': 'resale_customers',
18    'Average No. of Customers - All Other': 'other_customers'
19 }, inplace=True)
20
21 print(df_rename.columns)

Index(['RY', 'ToU', 'U', 'ORoRS', 'ORoCIS', 'ORoSR', 'ORoAOS', 'ASforR',
      'ASforCI', 'ASforSR', 'ASforAO', 'ANoCR', 'ANoCCI', 'ANoCSR', 'ANoCAO'],
      dtype='object')

1 # df_rename.columns

```

```

1 df_energy.columns

```

```
Index(['RY', 'ORoRS', 'ORoCIS', 'ORoSR', 'ORoAOS', 'ASforR', 'ASforCI',
      'ASforSR', 'ASforAO', 'ANoCR',
      ...
      'U_Westfield', 'U_Whittemore', 'U_Wilton', 'U_Wilton ', 'U_Winteraset',
      'U_Winteraset ', 'U_Woodbine', 'U_Woodbine ',
      'U_Woodbury County Rural Electric Cooperative', 'U_Woolstock'],
      dtype='object', length=239)
```

```
1 # Check Unique Values for each variable
2 def get_uniquevalues(df1):
3     unique_values=df1.apply(pd.Series.unique)
4     return unique_values
5
6 unq_values = get_uniquevalues(df)
7
8 for i in df.columns.tolist():
9     print("No. of unique values in ",i,"is",df[i].nunique())
```

```
No. of unique values in RY is 8
No. of unique values in ToU is 5
No. of unique values in U is 221
No. of unique values in ORoRS is 1446
No. of unique values in ORoCIS is 1422
No. of unique values in ORoSR is 302
No. of unique values in ORoAOS is 1251
No. of unique values in ASforR is 1398
No. of unique values in ASforCI is 1378
No. of unique values in ASforSR is 289
No. of unique values in ASforAO is 1006
No. of unique values in ANoCR is 987
No. of unique values in ANoCCI is 569
No. of unique values in ANoCSR is 14
No. of unique values in ANoCAO is 166
```


```
1 # Separate columns in list for better analysis
2 gen_cols=['reporting_year', 'utility_type', 'utility']
3 rev_cols=['residential_revenues', 'commercial_revenues', 'resale_revenues', 'other_revenues']
4 sal_cols=['residential_sales', 'commercial_sales', 'resale_sales', 'other_sales']
5 cus_cols=['residential_customers', 'commercial_customers', 'resale_customers', 'other_customers']
```



5. Data Vizualization


```
1 # Chart - 01 visualization
2 # Dependent variable "ORoCIS - commercial_revenues"
3 plt.figure(figsize=(5,5))
4 sns.distplot(df_energy['ANoCCI'], color = 'Blue')
```

```
<Axes: xlabel='ANoCCI', ylabel='Density'>
```



```
1 # Chart - 02 visualization
2 # Dependent variable "ASforR - residential_sales"
3 plt.figure(figsize=(5,5))
4 sns.distplot(df_energy['ASforR'], color = 'Blue')
```

```
<Axes: xlabel='ASforR', ylabel='Density'>
```




```

1 # Chart - 03 visualization
2 # Dependent variable "ANoCCI - commercial_customers"
3 plt.figure(figsize=(5,5))
4 sns.distplot(df_energy['ANoCCI'], color = 'Blue')

```

<Axes: xlabel='ANoCCI', ylabel='Density'>




```

1 # Display the heatmap
2 data['ToU'] = data['ToU'].astype('category').cat.codes
3 data['U'] = data['U'].astype('category').cat.codes
4
5 correlation_matrix = data.corr()
6
7 plt.figure(figsize=(12, 10))
8 sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', linewidths=0.5)
9
10 plt.title('Correlation Matrix Heatmap')
11 plt.show()

```

<Figure: Correlation Matrix Heatmap>





```

1 # Handling outliers & outlier treatments
2 df = df_energy.copy()
3 col_list = list(df.describe().columns)
4
5 # Find the outliers using boxplot
6 plt.figure(figsize=(25, 20))
7 plt.suptitle("Box Plot", fontsize=18, y=0.95)
8
9 for n, ticker in enumerate(col_list):
10
11     ax = plt.subplot(8, 4, n + 1)
12
13     plt.subplots_adjust(hspace=0.5, wspace=0.2)
14
15     sns.boxplot(x=df[ticker], color='pink', ax = ax)
16
17     ax.set_title(ticker.upper())

```



Box Plot



6. Feature Selection

```

1 # Feature Selection using PCA
2 from sklearn.preprocessing import StandardScaler
3 from sklearn.decomposition import PCA
4
5 df_energy = pd.DataFrame(data)
6
7 df_energy = pd.get_dummies(df_energy, columns=['ToU', 'U'])
8
9 scaler = StandardScaler()
10 scaled_features = scaler.fit_transform(df_energy)
11
12 # Set the number of principal components
13 pca = PCA(n_components=5)
14 principal_components = pca.fit_transform(scaled_features)
15
16 pca_df = pd.DataFrame(data=principal_components, columns=[f'PC{i+1}' for i in range(principal_components.shape[1])])
17
18 print("PCA result:")
19 print(pca_df)
20
21 print("Explained variance ratio by each principal component:")
22 print(pca.explained_variance_ratio_)

```



PCA result:

	PC1	PC2	PC3	PC4	PC5
0	-0.497902	-0.958827	0.137331	0.008330	0.714779
1	19.175408	1.009345	-3.442617	15.177474	0.370857
2	35.553614	-1.416309	0.115059	-8.228335	1.160257
3	1.813677	1.168169	-0.612939	3.710831	0.692447
4	-0.451013	-0.968069	0.105568	-0.059686	0.614929
...
1447	-0.159398	3.026250	-0.585653	-0.418187	-1.024038
1448	-0.293829	3.008457	-0.663307	-0.443012	-1.887233
1449	-0.347229	3.026585	-0.541275	-0.369408	-0.474269
1450	-0.151183	3.022306	-0.602712	-0.395926	-0.786458
1451	-0.246829	3.025089	-0.591248	-0.500880	-0.502602

[1452 rows x 5 columns]
Explained variance ratio by each principal component:
[0.04797756 0.01228517 0.00900024 0.00770409 0.00581886]

```

1 import pandas as pd
2
3 # Get the explained variance ratio of each principal component
4 explained_variance_ratio = pca.explained_variance_ratio_
5
6 # Create a DataFrame to store the results
7 pca_results = pd.DataFrame({'Principal Component': [f'PC{i+1}' for i in range(len(explained_variance_ratio))],
8                             'Explained Variance Ratio': explained_variance_ratio})
9
10 # Print the results
11 print(pca_results)

```

```

Principal Component  Explained Variance Ratio
0                   PC1                   0.047978
1                   PC2                   0.012285
2                   PC3                   0.009000
3                   PC4                   0.007704
4                   PC5                   0.005819

```

```

1 # Get the loadings of the principal components
2 df_raw = pd.read_csv('/content/drive/MyDrive/Machine Learning Project/electricity_consumption_data.csv')
3 pca = PCA(n_components=5)
4 pca.fit(df)
5 loadings = pca.components_
6
7 # Create a DataFrame to store the loadings
8 loadings_df = pd.DataFrame(data=loadings, columns=df.columns)
9
10 # Print the loadings
11 print(loadings_df)

```

```

RY      ORoRS      ORoCIS      ORoSR      ORoAOS      ASforR      ASforCI \
0  2.328384e-10  0.506407  0.844632  0.164063  0.055135  0.004044  0.011580
1  -2.947949e-09 -0.173898 -0.094576  0.964923  0.170622  0.001558  0.005645
2  -1.068387e-08  0.725235 -0.467292 -0.004816  0.504748  0.015913  0.020676
3  2.979599e-08 -0.432652  0.243207 -0.204232  0.841752  0.011515  0.040726
4  -2.953693e-07  0.003975 -0.010586 -0.012222 -0.065947  0.203717  0.574630

ASforSR      ASforAO      ANoCR      ...      U_Westfield      U_Whittemore \
0  0.006417  0.000624  0.000423  ... -6.314595e-12 -6.176206e-12
1  0.024266  0.001915  0.000084  ... -1.833664e-12 -3.454049e-12
2  0.012083  0.007012  0.001353  ... -2.169157e-10 -1.971119e-10
3  0.038901  0.010027  0.000958  ...  1.860925e-10  1.792167e-10
4  0.780912  0.116294  0.018016  ...  6.180027e-10  4.462237e-10

U_Wilton      U_Wilton      U_Winterset      U_Winterset      U_Woodbine \
0 -4.311260e-12 -1.431225e-12 -3.786683e-12 -1.267996e-12 -4.520485e-12
1 -6.495211e-12 -2.273773e-12 -1.387030e-11 -4.440536e-12  7.316648e-13
2 -8.680589e-11 -2.788794e-11  2.186350e-11 -1.421687e-11 -6.641466e-11
3  1.088306e-10  3.347575e-11  4.770152e-12  9.678523e-12  1.109183e-10
4  4.834784e-11  4.015610e-11  3.729772e-10  2.063949e-10 -4.351868e-10

U_Woodbine      U_Woodbury County Rural Electric Cooperative      U_Woolstock
0 -1.512446e-12                                     -4.317897e-12 -6.281196e-12
1 -1.154734e-12                                     -3.705361e-11 -2.099772e-12
2 -2.392460e-11                                     9.106545e-10 -2.272202e-10
3  4.756345e-11                                     -6.771194e-10  1.926730e-10
4 -1.551476e-10                                     3.080807e-09  5.603604e-10

[5 rows x 239 columns]

```

```

1 most_important_feature_pc1 = loadings_df.iloc[:, 0].abs().idxmax()
2 print(most_important_feature_pc1)

```

```
4
```

```

1 # Select PC1 as the feature
2 X = pca_df[['PC1']]
3
4 # Assuming ORoRS as the dependent variable for regression
5 y = df_energy['ORoRS']
6
7 # Split the dataset into training and testing sets
8 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
9
10 # Print the shapes of the resulting datasets
11 print("Shapes of the datasets:")
12 print(f"X_train: {X_train.shape}")
13 print(f"X_test: {X_test.shape}")
14 print(f"y_train: {y_train.shape}")
15 print(f"y_test: {y_test.shape}")

```

```

Shapes of the datasets:
X_train: (1161, 1)
X_test: (291, 1)
y_train: (1161,)
y_test: (291,)

```

7. Model Selection

```

1 from sklearn.linear_model import LinearRegression
2
3 # Assuming df_energy, pca_df, X, y, X_train, X_test, y_train, and y_test are already defined
4
5 # Initialize and train the Linear Regression model
6 linear_regressor = LinearRegression()
7 linear_regressor.fit(X_train, y_train)
8
9 # Predict on the test set
10 y_pred = linear_regressor.predict(X_test)
11
12 # Evaluate the model
13 mse = mean_squared_error(y_test, y_pred)
14 rmse = np.sqrt(mse)
15 r2 = r2_score(y_test, y_pred)
16
17 print("Linear Regression Model Evaluation:")


```

```

18 print(f"Mean Squared Error (MSE): {mse}")
19 print(f"Root Mean Squared Error (RMSE): {rmse}")
20 print(f"R-squared (R2): {r2}")
21
22 # Plotting the results
23 plt.scatter(X_test, y_test, color='blue', label='Actual')
24 plt.scatter(X_test, y_pred, color='red', label='Predicted')
25 plt.xlabel('PC1')
26 plt.ylabel('ORoRS')
27 plt.title('Linear Regression: Actual vs Predicted')
28 plt.legend()
29 plt.show()
30

```

Linear Regression Model Evaluation:
Mean Squared Error (MSE): 231588767751297.16
Root Mean Squared Error (RMSE): 15218040.864424605
R-squared (R2): 0.9067015636693613



```

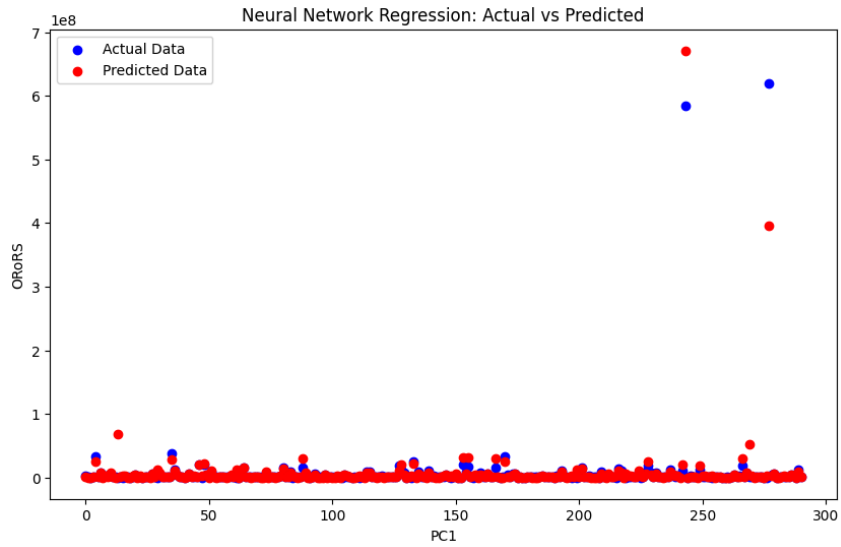
1 import tensorflow as tf
2 from tensorflow import keras
3 from tensorflow.keras import layers
4
5 # Separate features and target
6 features = df_energy.drop(columns=['ORoRS'])
7 target = df_energy['ORoRS']
8
9 # Standardize the features and target separately
10 scaler_features = StandardScaler()
11 scaled_features = scaler_features.fit_transform(features)
12
13 scaler_target = StandardScaler()
14 scaled_target = scaler_target.fit_transform(target.values.reshape(-1, 1))
15
16 # Select PC1 as the feature
17 X = pca_df[['PC1']]
18
19 # Use the scaled target for regression
20 y = scaled_target
21
22 # Split the dataset into training and testing sets
23 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
24
25 # Define the neural network model
26 model = keras.Sequential([
27     layers.Input(shape=(X_train.shape[1],)), # Input layer with the number of PCs as input shape
28     layers.Dense(32, activation='relu'), # Hidden layer with 32 neurons and ReLU activation
29     layers.Dense(1) # Output layer with a single neuron (for regression)
30 ])
31
32 model.compile(optimizer='adam', loss='mean_squared_error')
33
34 # Train the model
35 history = model.fit(X_train, y_train, epochs=100, batch_size=32, validation_data=(X_test, y_test), verbose=0)
36
37 # Evaluate the model on the test data
38 test_loss = model.evaluate(X_test, y_test)
39 print(f"Test Loss: {test_loss:.4f}")
40
41 # Make predictions on the test data
42 y_pred = model.predict(X_test)
43
44 # Inverse transform predictions and true values
45 y_pred_inv = scaler_target.inverse_transform(y_pred)
46 y_test_inv = scaler_target.inverse_transform(y_test)
47
48 # Evaluate the model
49 mse = mean_squared_error(y_test_inv, y_pred_inv)
50 rmse = np.sqrt(mse)
51 r2 = r2_score(y_test_inv, y_pred_inv)
52 print("Neural Network Regression Model Evaluation:")
53 print(f'Mean Squared Error (MSE): {mse}')
54 print(f'Root Mean Squared Error (RMSE): {rmse}')
55 print(f'R-squared (R2): {r2}')
56
57 # Plot the actual data and model predictions
58 plt.figure(figsize=(10, 6))
59 plt.scatter(range(len(y_test_inv)), y_test_inv, label='Actual Data', color='blue')
60 plt.scatter(range(len(y_pred_inv)), y_pred_inv, label='Predicted Data', color='red')
61 plt.xlabel('PC1')
62 plt.ylabel('ORoRS')
63 plt.legend()
64 plt.title('Neural Network Regression: Actual vs Predicted')
65 plt.show()

```

```

10/10 [=====] - 0s 3ms/step - loss: 0.0589
Test Loss: 0.0589
10/10 [=====] - 0s 2ms/step
Neural Network Regression Model Evaluation:
Mean Squared Error (MSE): 230707569326548.2
Root Mean Squared Error (RMSE): 15189060.844125558
R-squared (R2): 0.9070565654940369

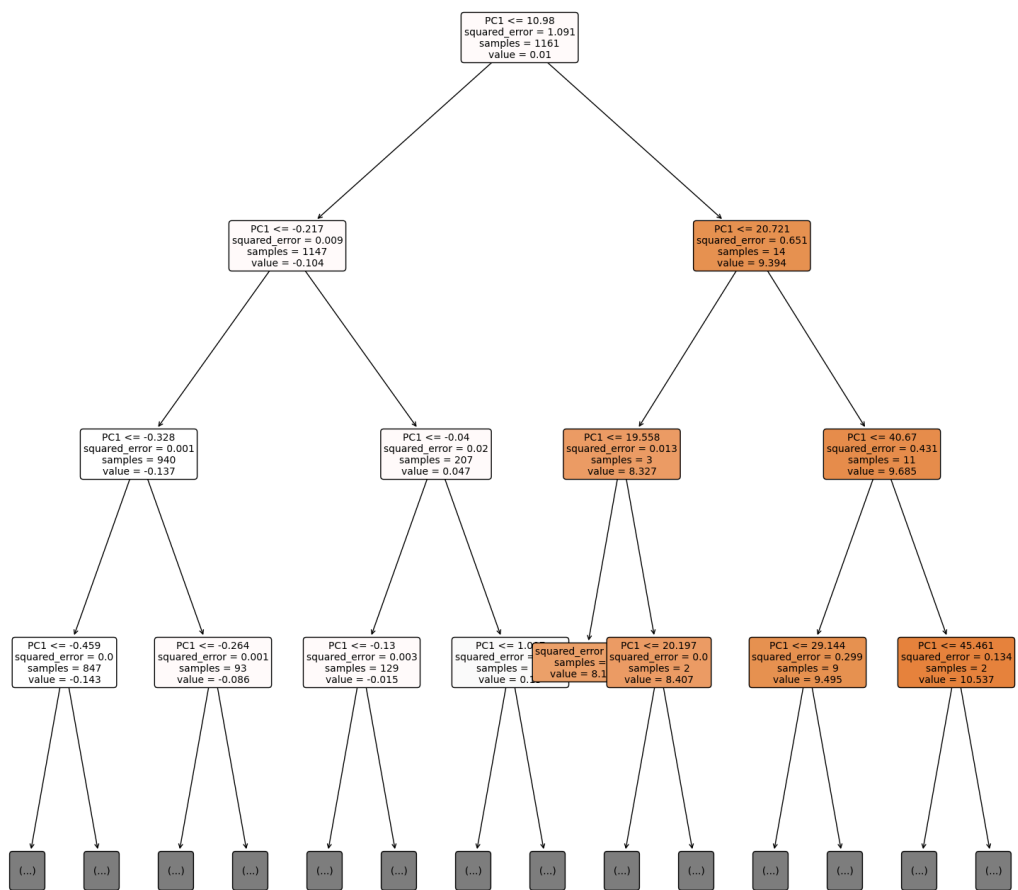
```



```

1 from sklearn.tree import DecisionTreeRegressor, plot_tree
2
3 # Create a decision tree regressor
4 regressor = DecisionTreeRegressor(random_state=42)
5
6 # Fit the regressor to the training data
7 regressor.fit(X_train, y_train)
8
9 # Visualize the decision tree
10 fig, ax = plt.subplots(figsize=(15, 15))
11 plot_tree(regressor, max_depth=3, feature_names=['PC1'], class_names=['ORoRS'],
12           filled=True, rounded=True, fontsize=10, label='all', ax=ax)
13 plt.tight_layout() # Adjust layout to prevent overlapping
14 plt.show()
15
16 # Make predictions on the testing data
17 y_pred = regressor.predict(X_test)
18
19 print("Decision Tree Model Evaluation:")
20
21 # Calculate the mean squared error
22 mse = mean_squared_error(y_test, y_pred)
23 print(f"Mean Squared Error (MSE): {mse}")
24
25 # Calculate the Root Mean Squared Error
26 rmse = np.sqrt(mse)
27 print(f"Root Mean Squared Error (RMSE): {rmse}")
28
29 # Calculate the R-squared score
30 r2 = r2_score(y_test, y_pred)
31 print(f"R-squared (R2): {r2}")

```



Decision Tree Model Evaluation:
Mean Squared Error (MSE): 0.001635622501685365
Root Mean Squared Error (RMSE): 0.04044283004050737
R-squared (R2): 0.997416876336296

```

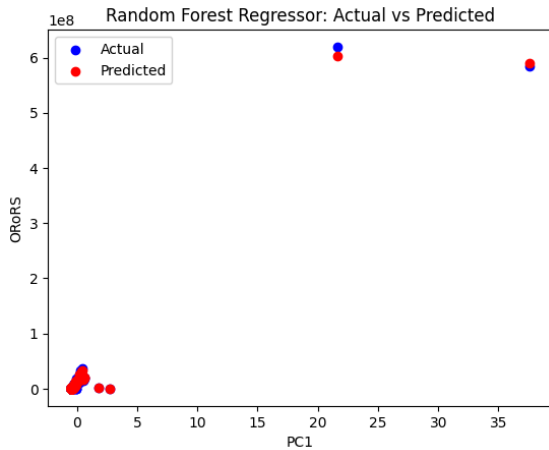
1 from sklearn.ensemble import RandomForestRegressor
2
3 # Select PC1 as the feature
4 X = pca_df[['PC1']]
5
6 # Assuming ORoRS as the dependent variable for regression
7 y = df_energy['ORoRS']
8
9 # Split the dataset into training and testing sets
10 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
11
12 # Initialize and train the Random Forest Regressor
13 rf_regressor = RandomForestRegressor(n_estimators=100, random_state=42)
14 rf_regressor.fit(X_train, y_train)
15
16 # Predict on the test set
17 y_pred = rf_regressor.predict(X_test)
18
19 # Evaluate the model
20 mse = mean_squared_error(y_test, y_pred)
21 rmse = np.sqrt(mse)
22 r2 = r2_score(y_test, y_pred)
23
24 print("Random Forest Regressor Model Evaluation:")
25 print(f"Mean Squared Error (MSE): {mse}")
26 print(f"Root Mean Squared Error (RMSE): {rmse}")
27 print(f"R-squared (R2): {r2}")
28
29 # Plotting the results
30 import matplotlib.pyplot as plt
31
32 plt.scatter(X_test, y_test, color='blue', label='Actual')
33 plt.scatter(X_test, y_pred, color='red', label='Predicted')
34 plt.xlabel('PC1')
35 plt.ylabel('ORoRS')
36 plt.title('Random Forest Regressor: Actual vs Predicted')
37 plt.legend()
38 plt.show()

```

```

Random Forest Regressor Model Evaluation:
Mean Squared Error (MSE): 5297060534684.081
Root Mean Squared Error (RMSE): 2301534.387030548
R-squared (R2): 0.9978660128043624

```



```

1 import xgboost as xgb
2
3 # Assuming df_energy, pca_df, X, y, X_train, X_test, y_train, and y_test are already defined
4
5 # Initialize and train the XGBoost regression model
6 xgbr = xgb.XGBRegressor(verbosity=0)
7 xgbr.fit(X_train, y_train)
8
9 # Predictions on the test set
10 y_pred = xgbr.predict(X_test)
11
12 # Evaluate the model
13 mse = mean_squared_error(y_test, y_pred)
14 rmse = np.sqrt(mse)
15 r2 = r2_score(y_test, y_pred)
16
17 print("XGBoost Regression Model Evaluation:")
18 print(f"Mean Squared Error (MSE): {mse}")
19 print(f"Root Mean Squared Error (RMSE): {rmse}")
20 print(f"R-squared (R2): {r2}")
21
22 # Plotting the results
23 plt.scatter(X_test, y_test, color='blue', label='Actual')
24 plt.scatter(X_test, y_pred, color='red', label='Predicted')
25 plt.xlabel('PC1')
26 plt.ylabel('ORoRS')
27 plt.title('XGBoost Regression: Actual vs Predicted')
28 plt.legend()
29 plt.show()

```

XGBoost Regression Model Evaluation:
Mean Squared Error (MSE): 13617885819724.05
Root Mean Squared Error (RMSE): 3690241.9730586843
R-squared (R2): 0.9945138640986516



8. Model Evaluation

Based on the evaluation metrics from above five models, the Random Forest Regressor model demonstrates superior performance compared to the other models. It achieves this by exhibiting the lowest Mean Squared Error (MSE) and the highest R-squared value among all models. These metrics indicate that the Random Forest Regressor provides more accurate predictions and better explains the variance in the target variable compared to the other regression models.

```
1 # Evaluation results for each model
2 models = ['Linear Regression', 'Neural Network Regression', 'Decision Tree', 'Random Forest Regressor', 'XGBoost Regression']
3 mse_values = [231588760521001.2, 222835706743011.4, 6411923390724.246, 5297160256238.563, 13617885819724.05]
4 rmse_values = [15218040.626867875, 14927682.564383911, 2532177.5985748405, 2301556.051074699, 3690241.9730586843]
5 r2_values = [0.9067015665821766, 0.9102278439510403, 0.997416876336296, 0.9978659726302849, 0.9945138640986516]
6
7 # Plotting
8 fig, axs = plt.subplots(3, figsize=(15, 15))
9
10 # MSE comparison
11 axs[0].bar(models, mse_values, color=['blue', 'orange', 'green', 'red', 'purple'])
12 axs[0].set_title('Mean Squared Error (MSE) Comparison')
13 axs[0].set_ylabel('MSE')
14
15 # RMSE comparison
16 axs[1].bar(models, rmse_values, color=['blue', 'orange', 'green', 'red', 'purple'])
17 axs[1].set_title('Root Mean Squared Error (RMSE) Comparison')
18 axs[1].set_ylabel('RMSE')
19
20 # R-squared comparison
21 axs[2].bar(models, r2_values, color=['blue', 'orange', 'green', 'red', 'purple'])
22 axs[2].set_title('R-squared (R2) Comparison')
23 axs[2].set_ylabel('R-squared')
24
25 plt.tight_layout()
26 plt.show()
```

